

Introduction to TensorRT

Lijian Chen
12 Apr 2019

Overview

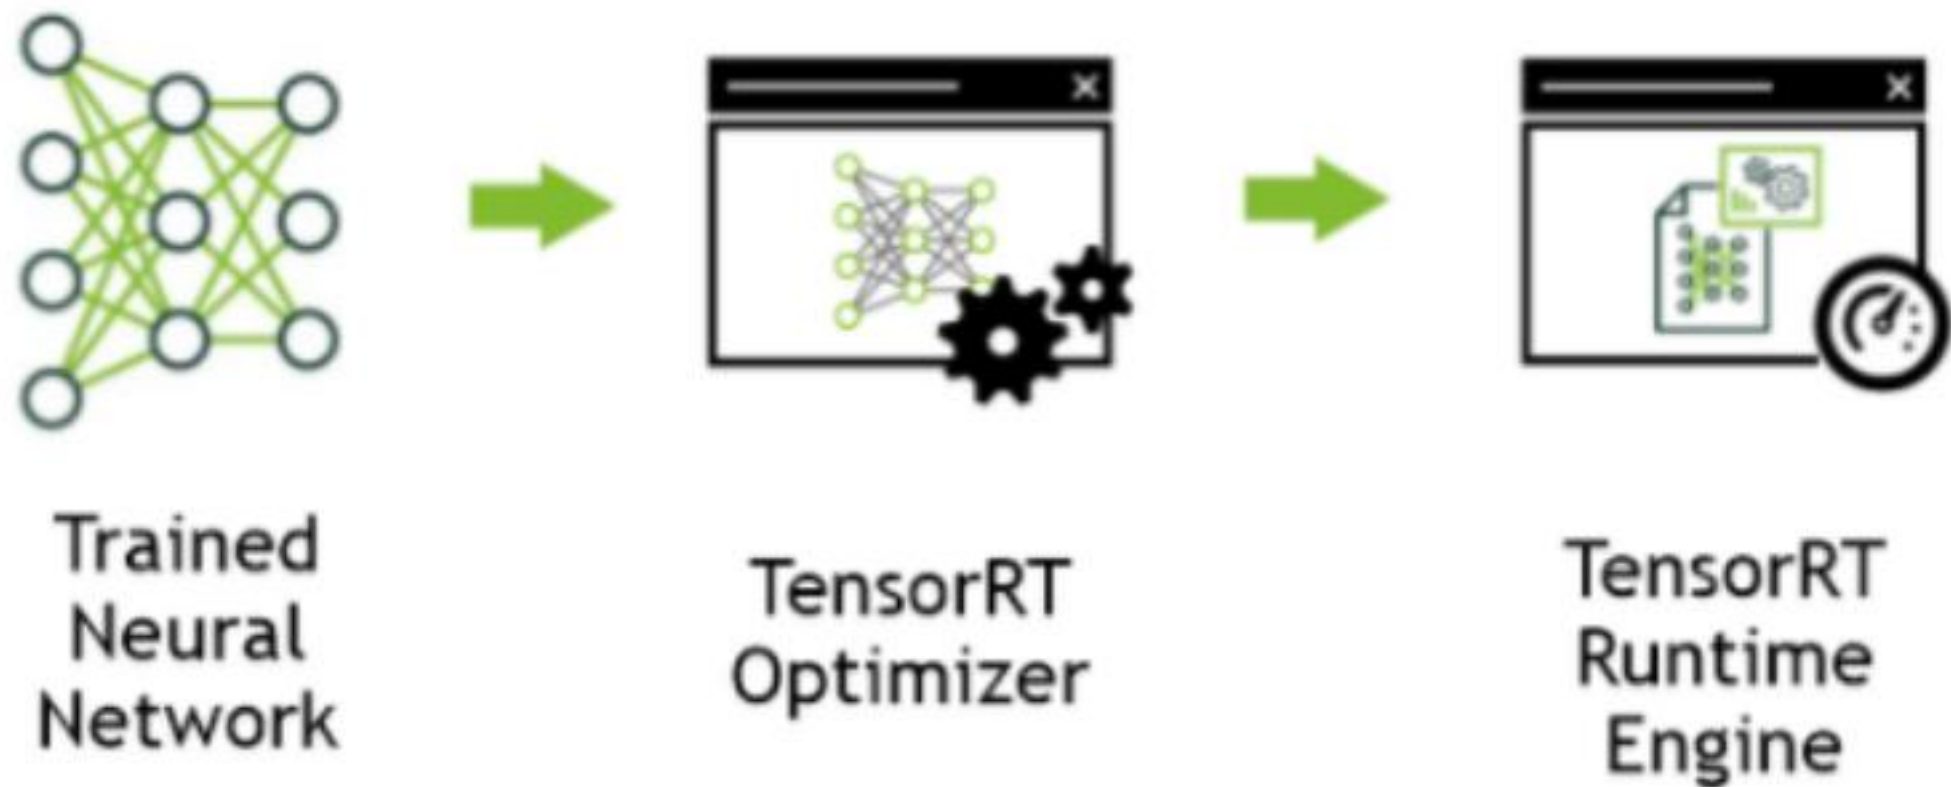
- ❖ How to measure deep learning inference ?
- ❖ What is TensorRT ?
- ❖ How does it work ?
- ❖ How to use TensorRT ?
- ❖ How to use TensorRT Inference Server ?

Measure Inference



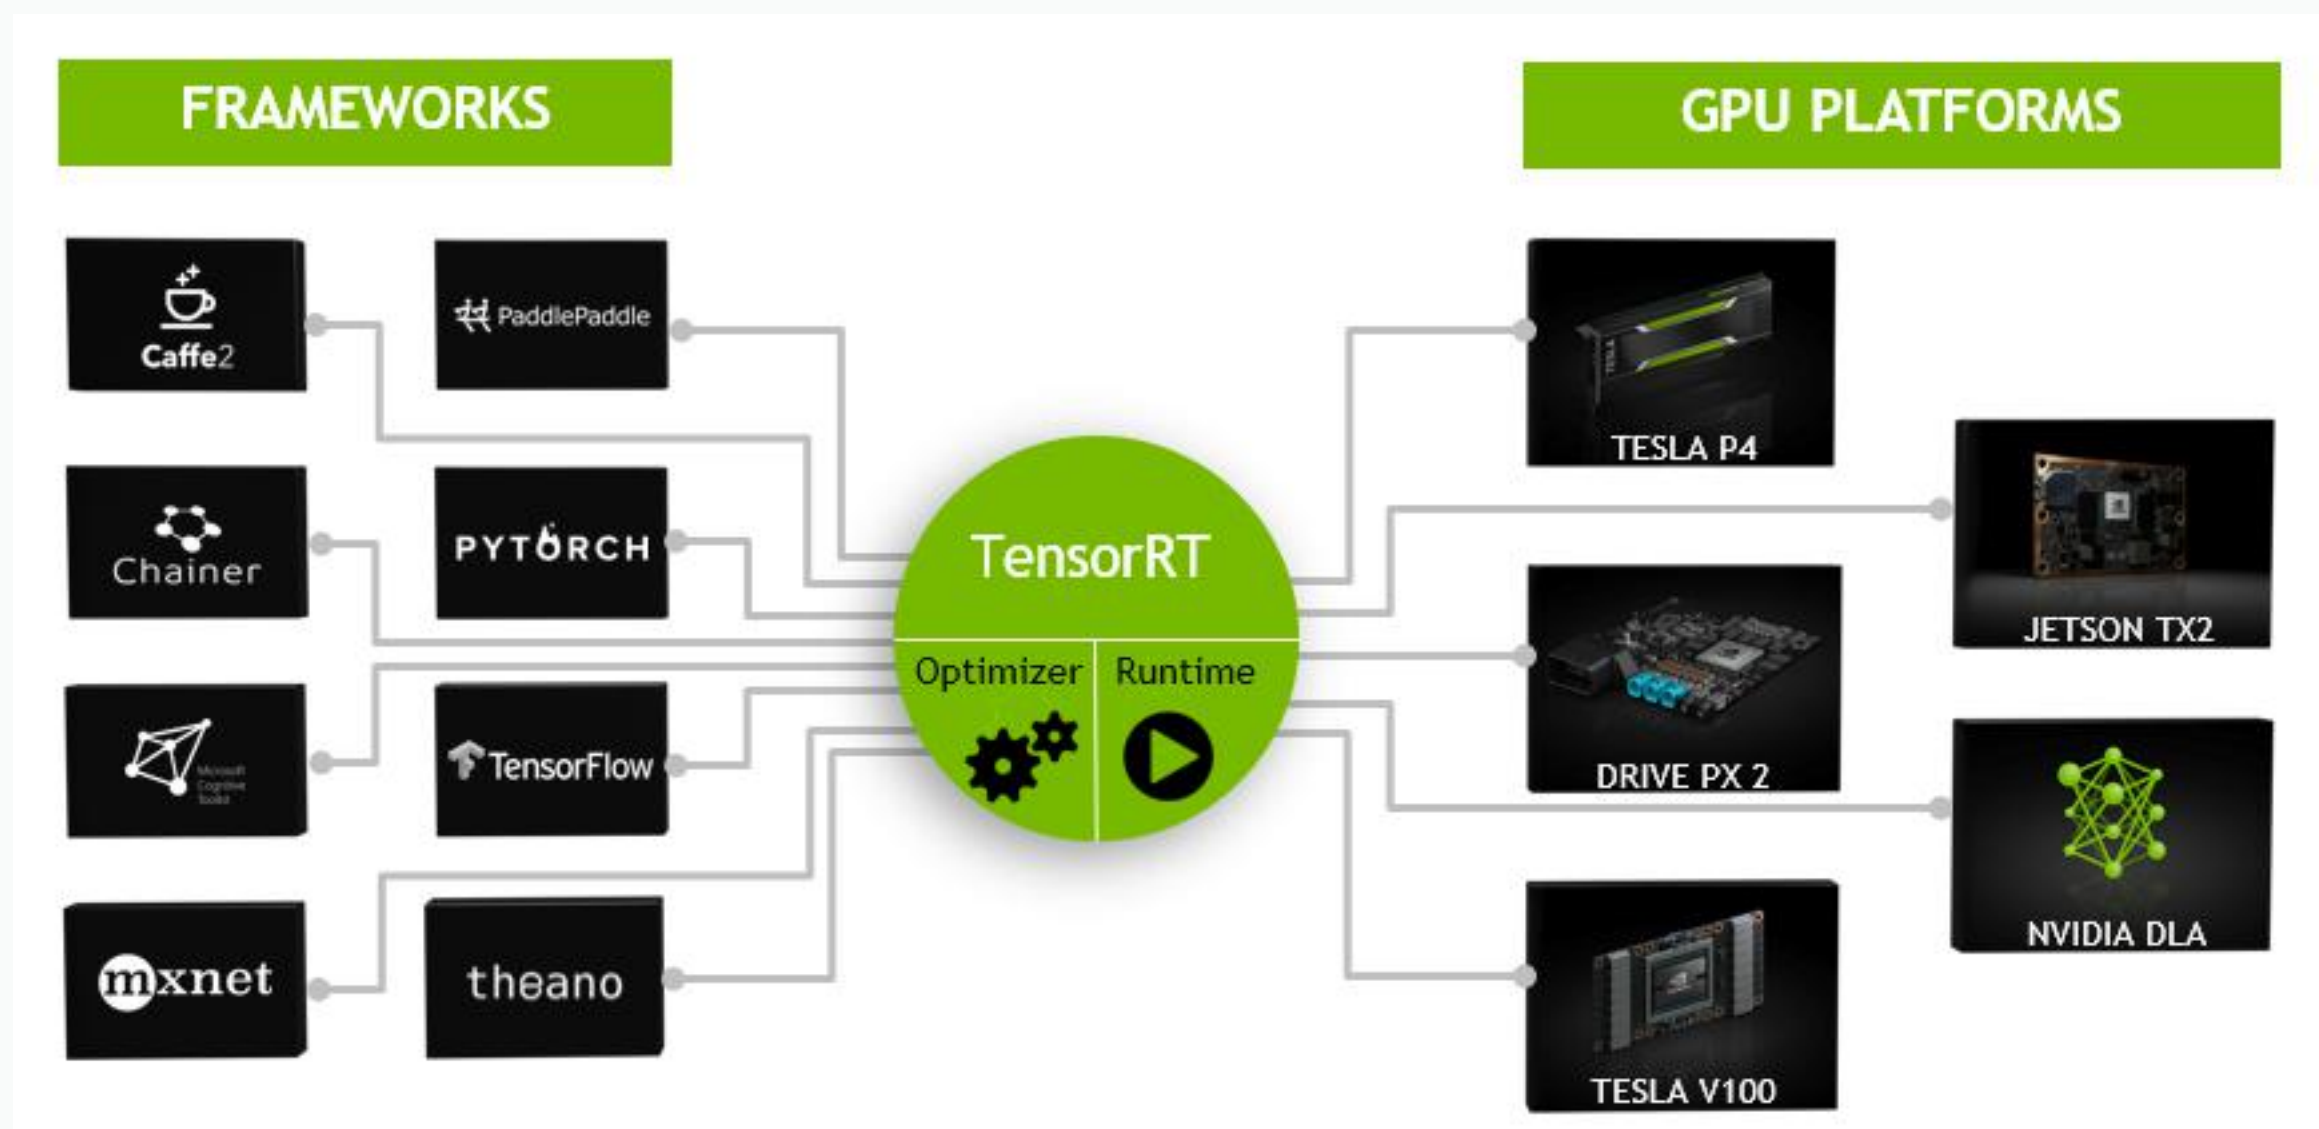
- ❖ Accuracy: ROC, Precision-Recall, confusion matrix, ...
- ❖ Latency: milliseconds / inference
- ❖ Throughput: inferences/second
- ❖ Memory usage: memory that need to be reserved to do inference
- ❖ Efficiency: throughput/watt

What is TensorRT



a C++ library
an optimiser + a runtime
provides C++ and Python API
integrated with frameworks

deep learning inference platform
support INT8, FP16, FP32 precisions.
GPU-specific



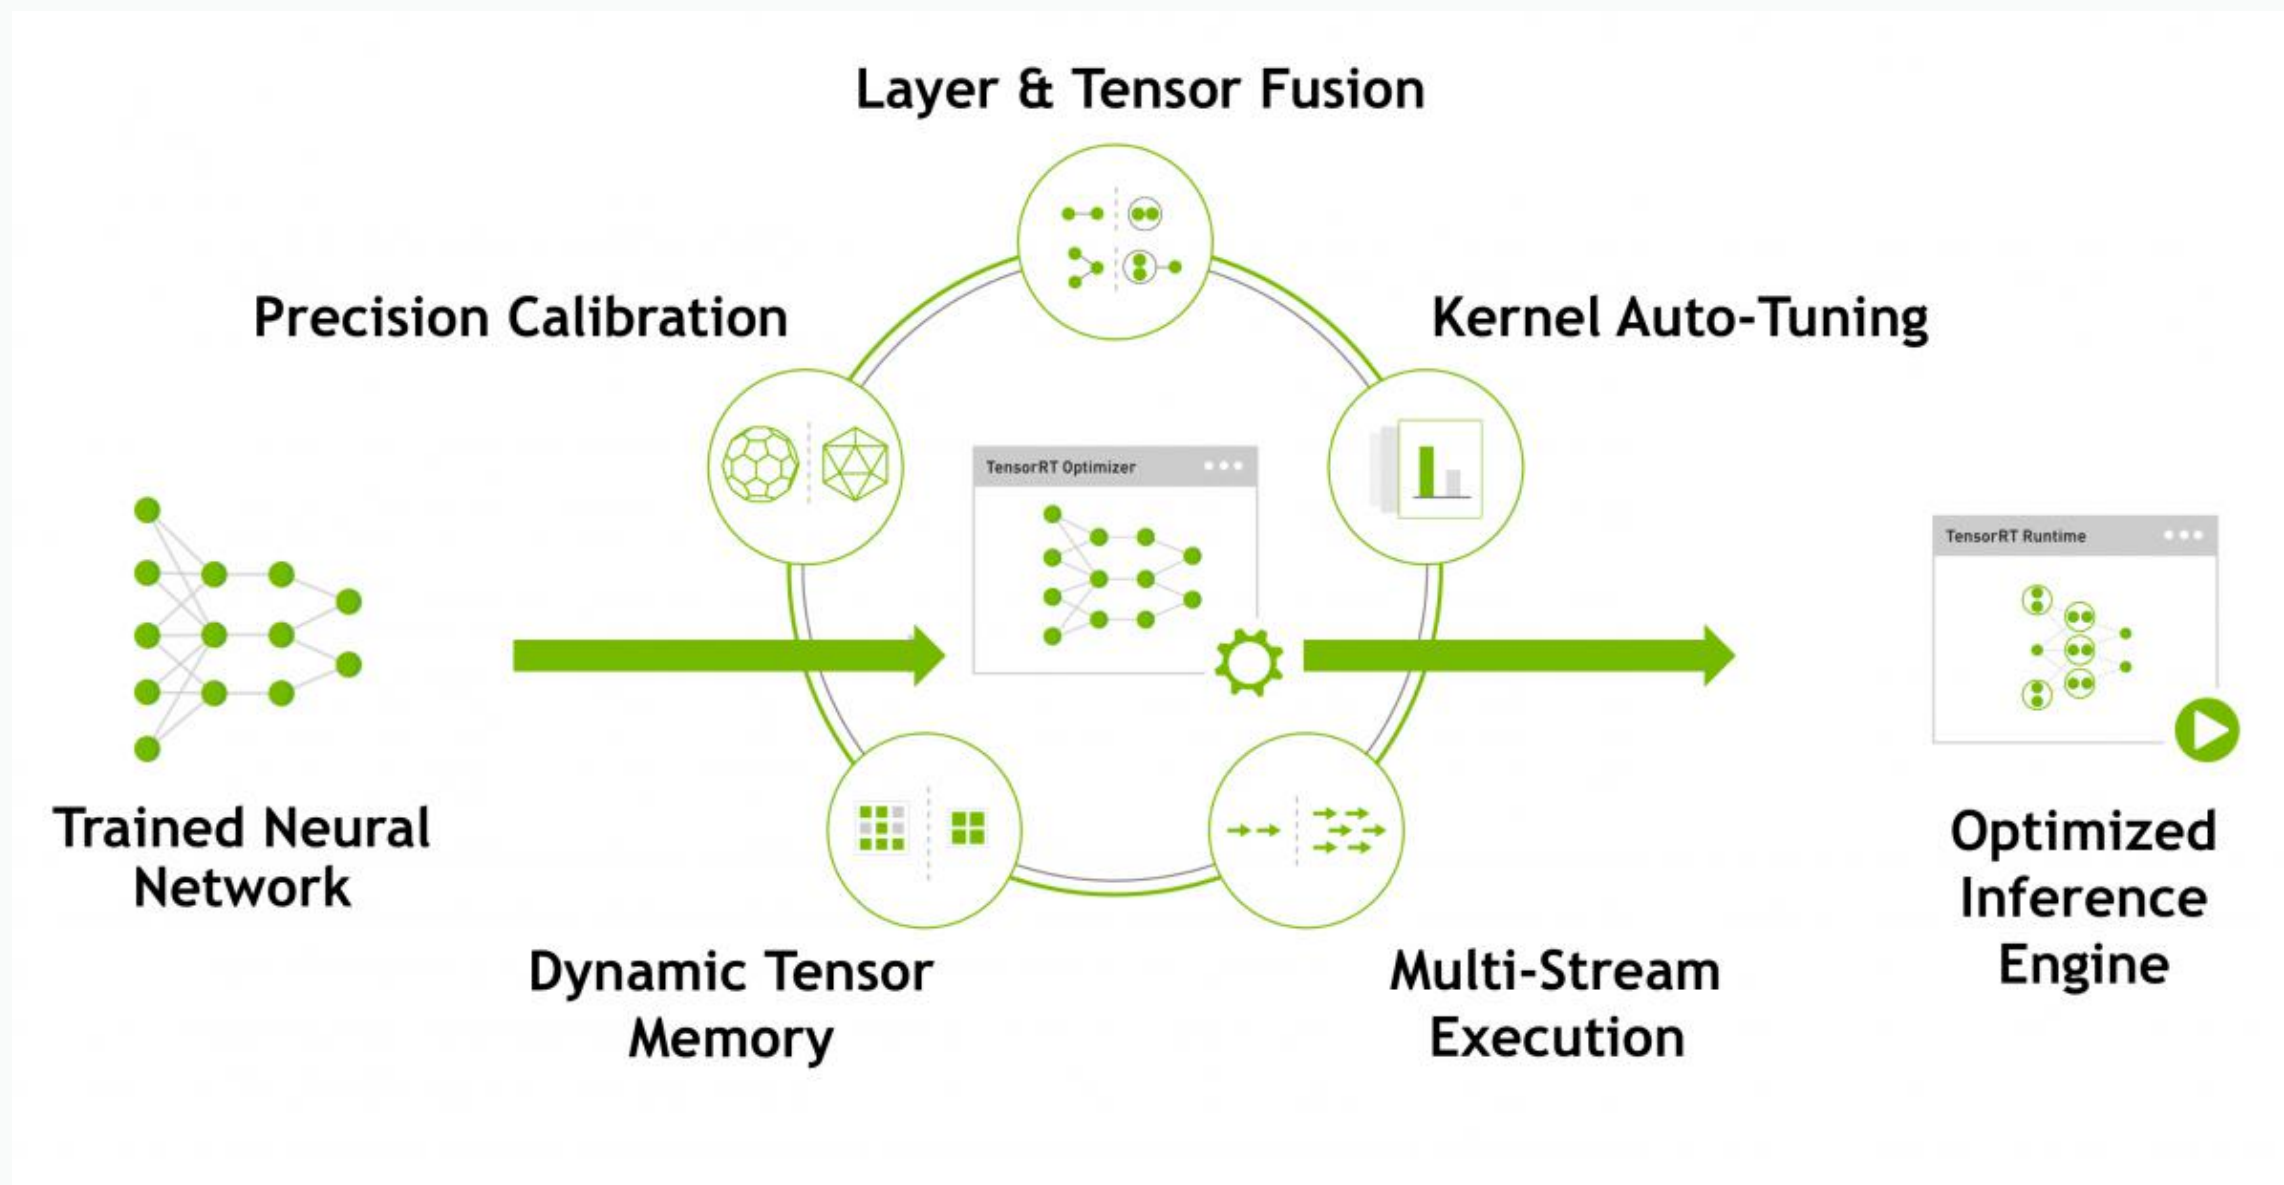
Widely Adopted



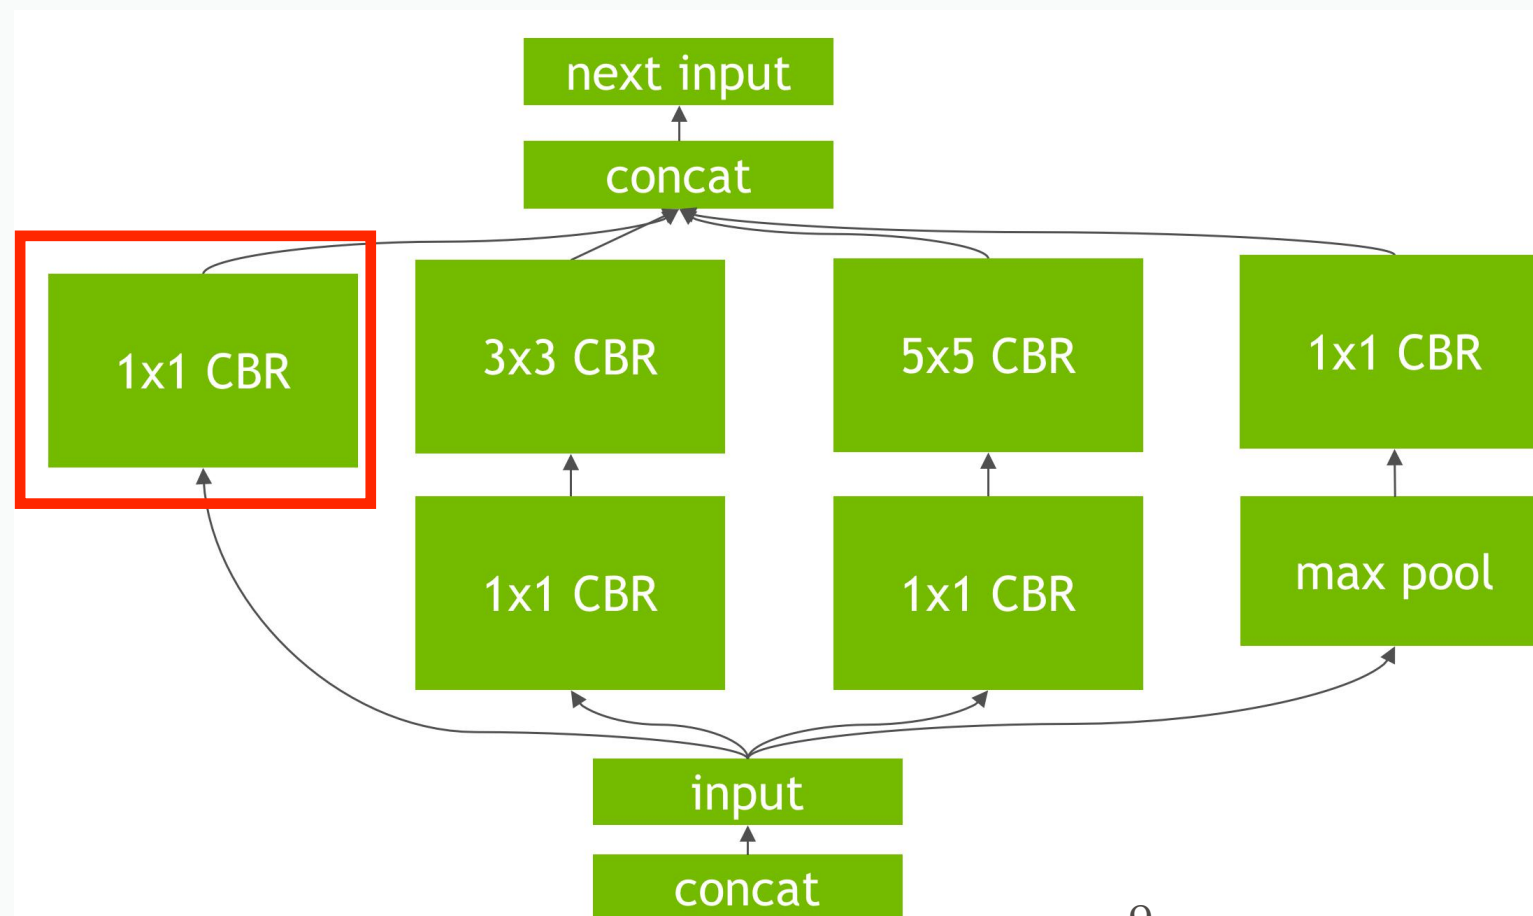
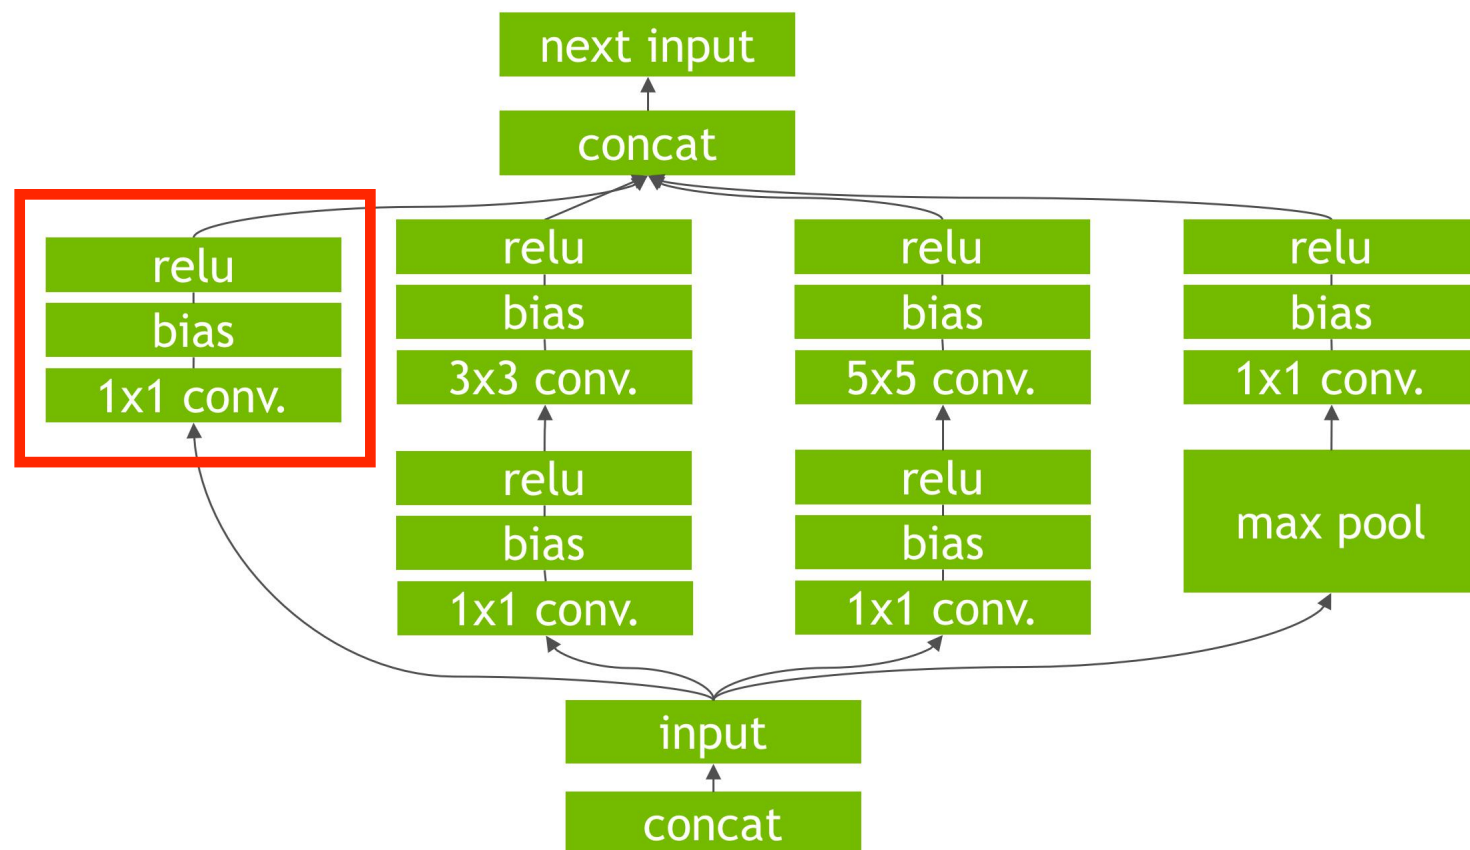
AURORA



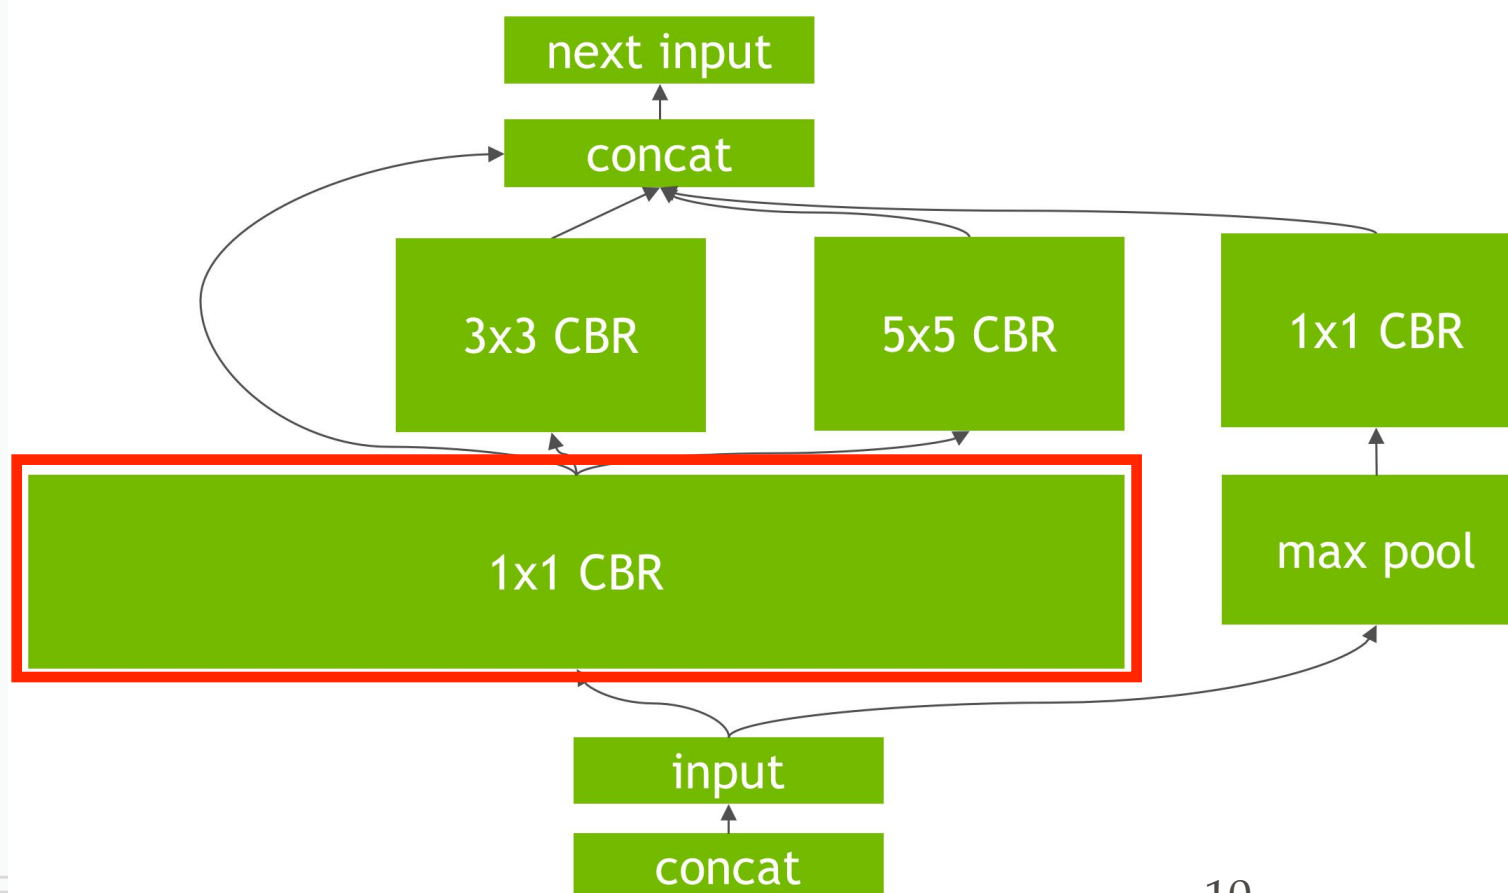
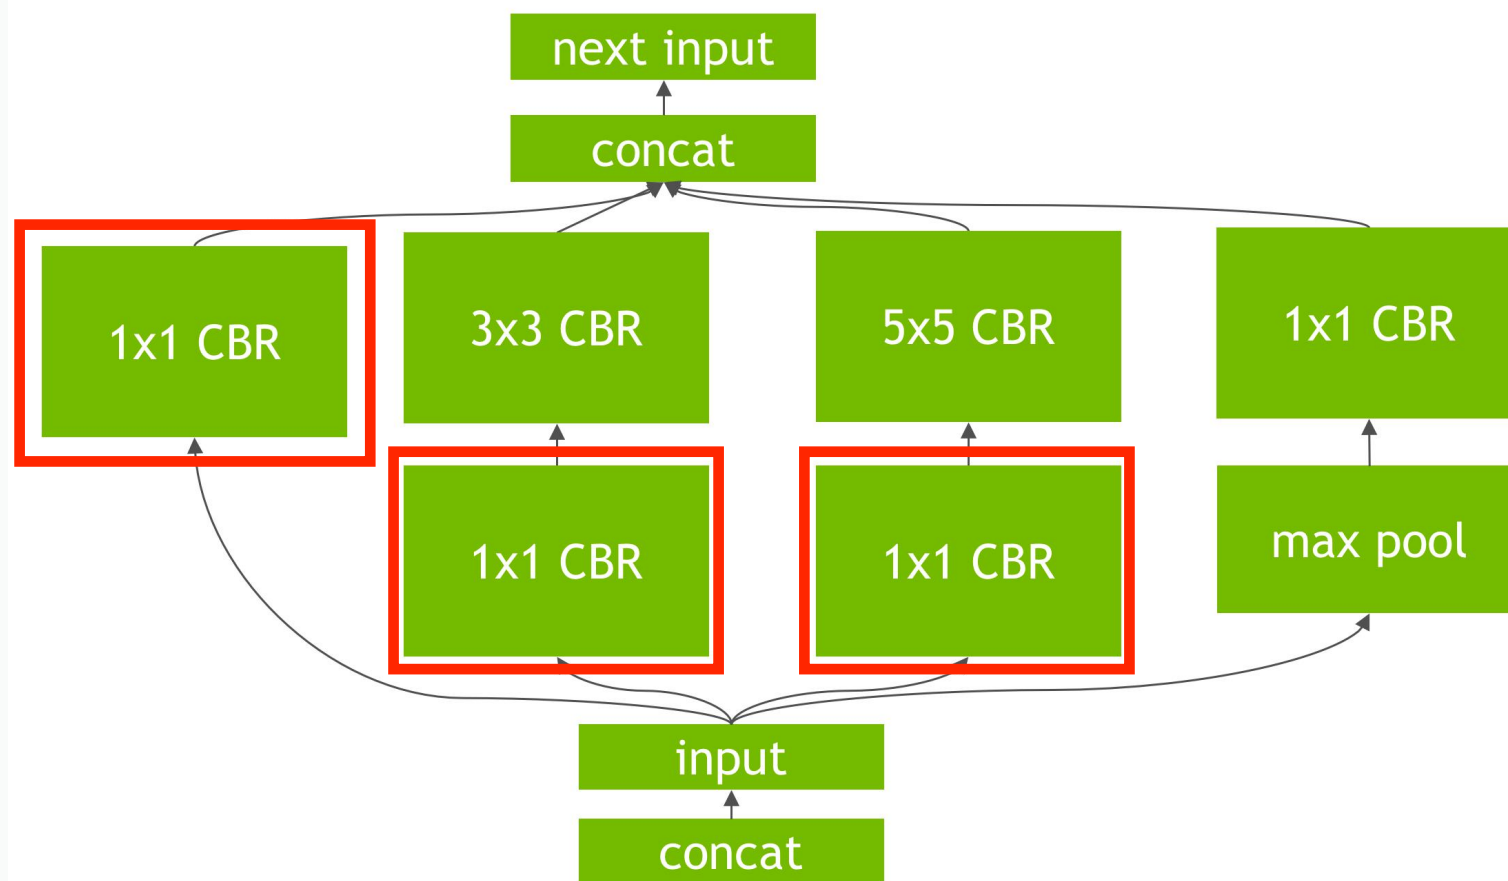
How it works



layer fusion example

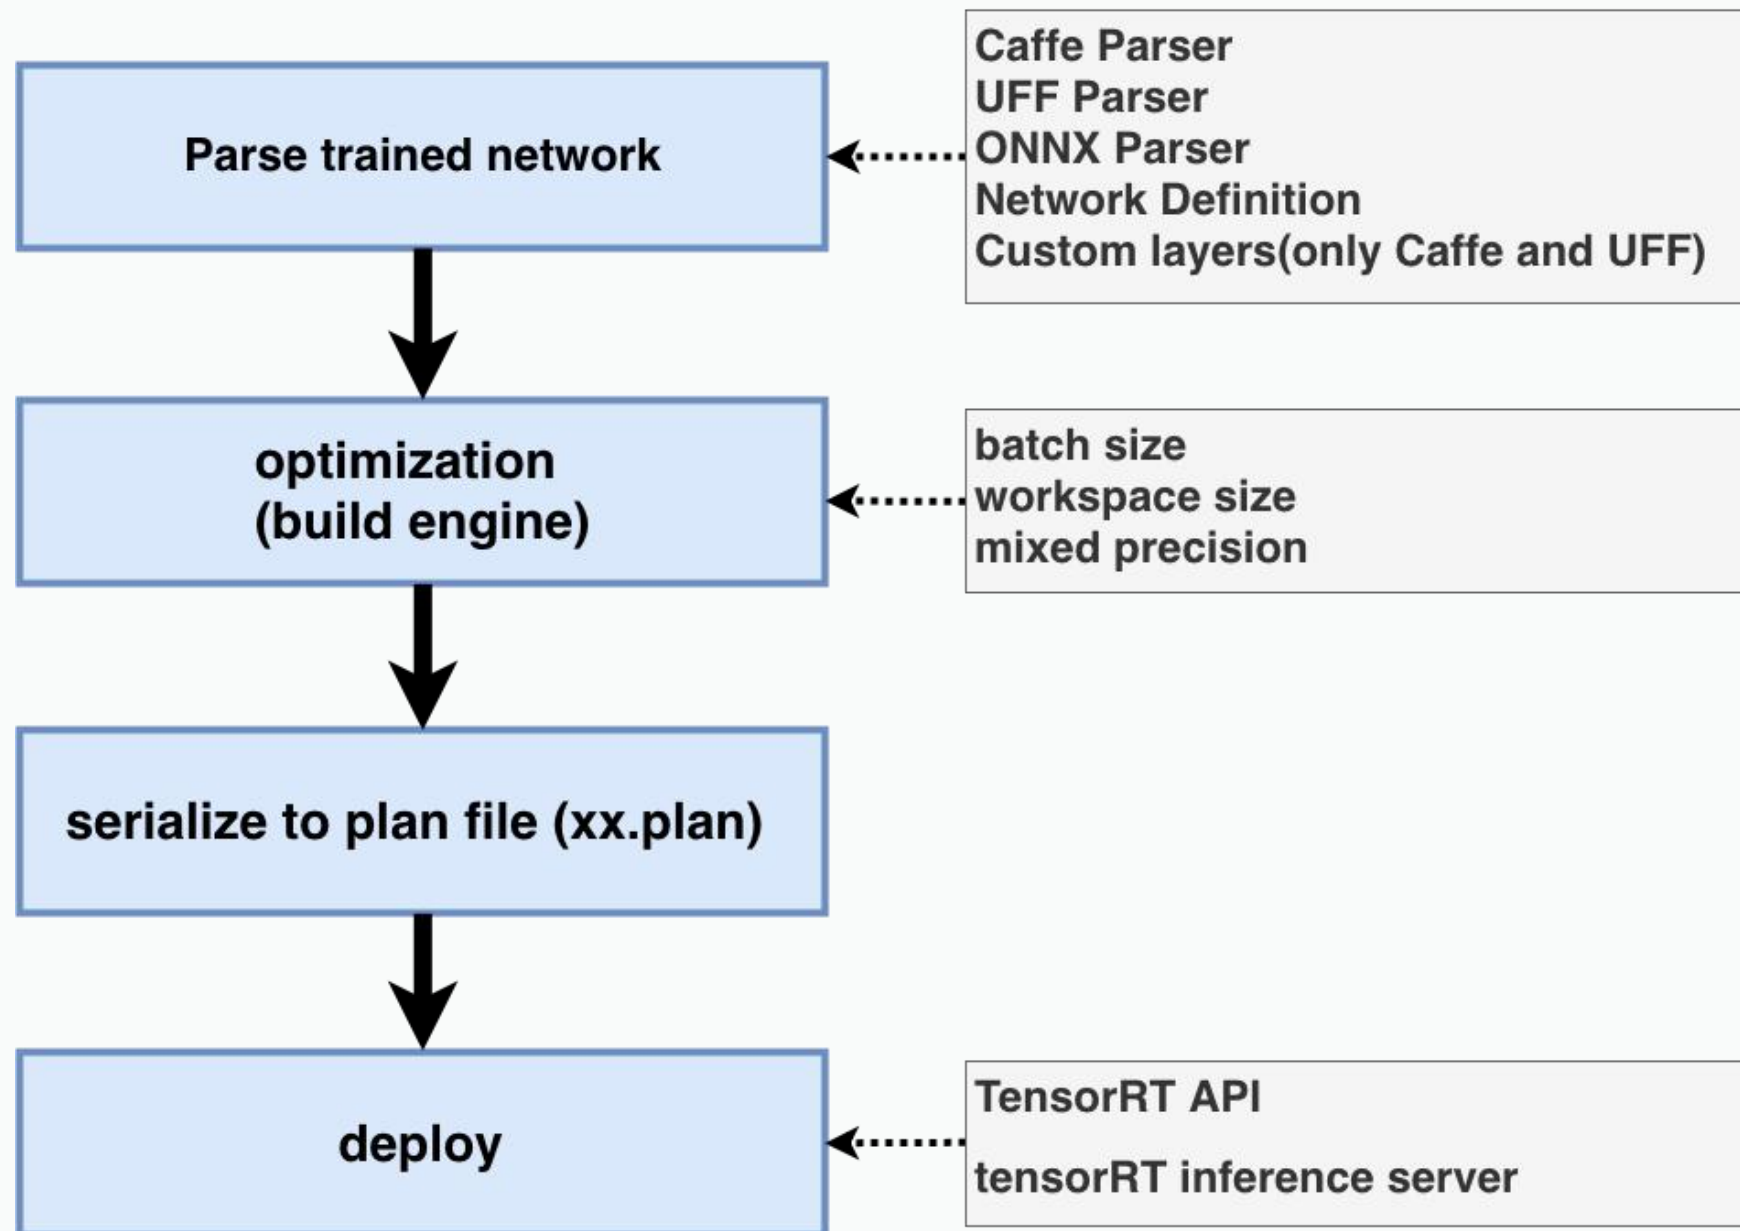


vertical fusion



horizontal fusion

How to use it



TensorRT Supported Layers

Layers

IConvolutionLayer
IFullyConnectedLayer
IActivationLayer
IPoolingLayer
ILRNLayer
IScaleLayer
ISoftMaxLayer
IConcatenationLayer
IDeconvolutionLayer
IElementWiseLayer
IGatherLayer

RNN Layers

IRNNLayer
IRNNv2Layer

IPluginLayer
IPluginV2Layer
IUnaryLayer
IReduceLayer
IPaddingLayer
IShuffleLayer
ISliceLayer
ITopKLayer
IMatrixMultiplyLayer
IRaggedSoftMaxLayer
IIdentityLayer
IConstantLayer

- RPROI_TRT
- Normalize_TRT
- PriorBox_TRT
- GridAnchor_TRT
- NMS_TRT
- LReLU_TRT
- Reorg_TRT
- Region_TRT
- Clip_TRT

Supported Layers in Frameworks

Caffe

These are the operations that are supported in a Caffe framework:

- BatchNormalization
- BNLL
- Clip⁵
- Concatenation
- Convolution
- Crop
- Deconvolution
- Dropout
- ElementWise
- ELU
- InnerProduct
- LeakyReLU
- LRN
- Permute
- Pooling
- Power
- Reduction
- ReLU, TanH, and Sigmoid
- Reshape
- SoftMax
- Scale

TensorFlow

These are the operations that are supported in a TensorFlow framework:

- Add, Sub, Mul, Div, Minimum and Maximum
- ArgMax
- ArgMin
- AvgPool
- BiasAdd
- Clip
- ConcatV2
- Const
- Conv2D
- ConvTranspose2D
- DepthwiseConv2dNative
- Elu
- ExpandDims
- FusedBatchNorm
- Identity
- LeakyReLU
- MaxPool
- Mean
- Negative, Abs, Sqrt, Recip, Rsqrt, Pow, Exp and Log
- Pad is supported if followed by one of these TensorFlow layers: Conv2D, DepthwiseConv2dNative, MaxPool, and AvgPool.
- Placeholder
- ReLU, TanH, and Sigmoid
- Relu6
- Reshape
- Sin, Cos, Tan, Asin, Acos, Atan, Sinh, Cosh, Asinh, Acosh, Atanh, Ceil and Floor
- Selu
- Slice
- SoftMax

Note: If the input to a TensorFlow `SoftMax` op is not `NHWC`, TensorFlow will automatically insert a transpose layer with a non-constant permutation, causing the UFF converter to fail. It is therefore advisable to manually transpose `SoftMax` inputs to `NHWC` using a constant permutation.

- Softplus
- Softsign
- Transpose

ONNX

Since the ONNX parser is an open source p

These are the operations that are supporte

- Abs
- Add
- ArgMax
- ArgMin
- AveragePool
- BatchNormalization
- Cast
- Ceil
- Clip
- Concat
- Constant
- Conv
- ConvTranspose
- DepthToSpace
- Div
- Dropout
- Elu
- Exp
- Flatten
- Floor
- Gather
- Gemm
- GlobalAveragePool
- GlobalMaxPool
- HardSigmoid
- Identity
- ImageScaler
- InstanceNormalization
- LRN
- LeakyRelU
- Log
- LogSoftmax
- MatMul
- Max
- MaxPool
- Mean

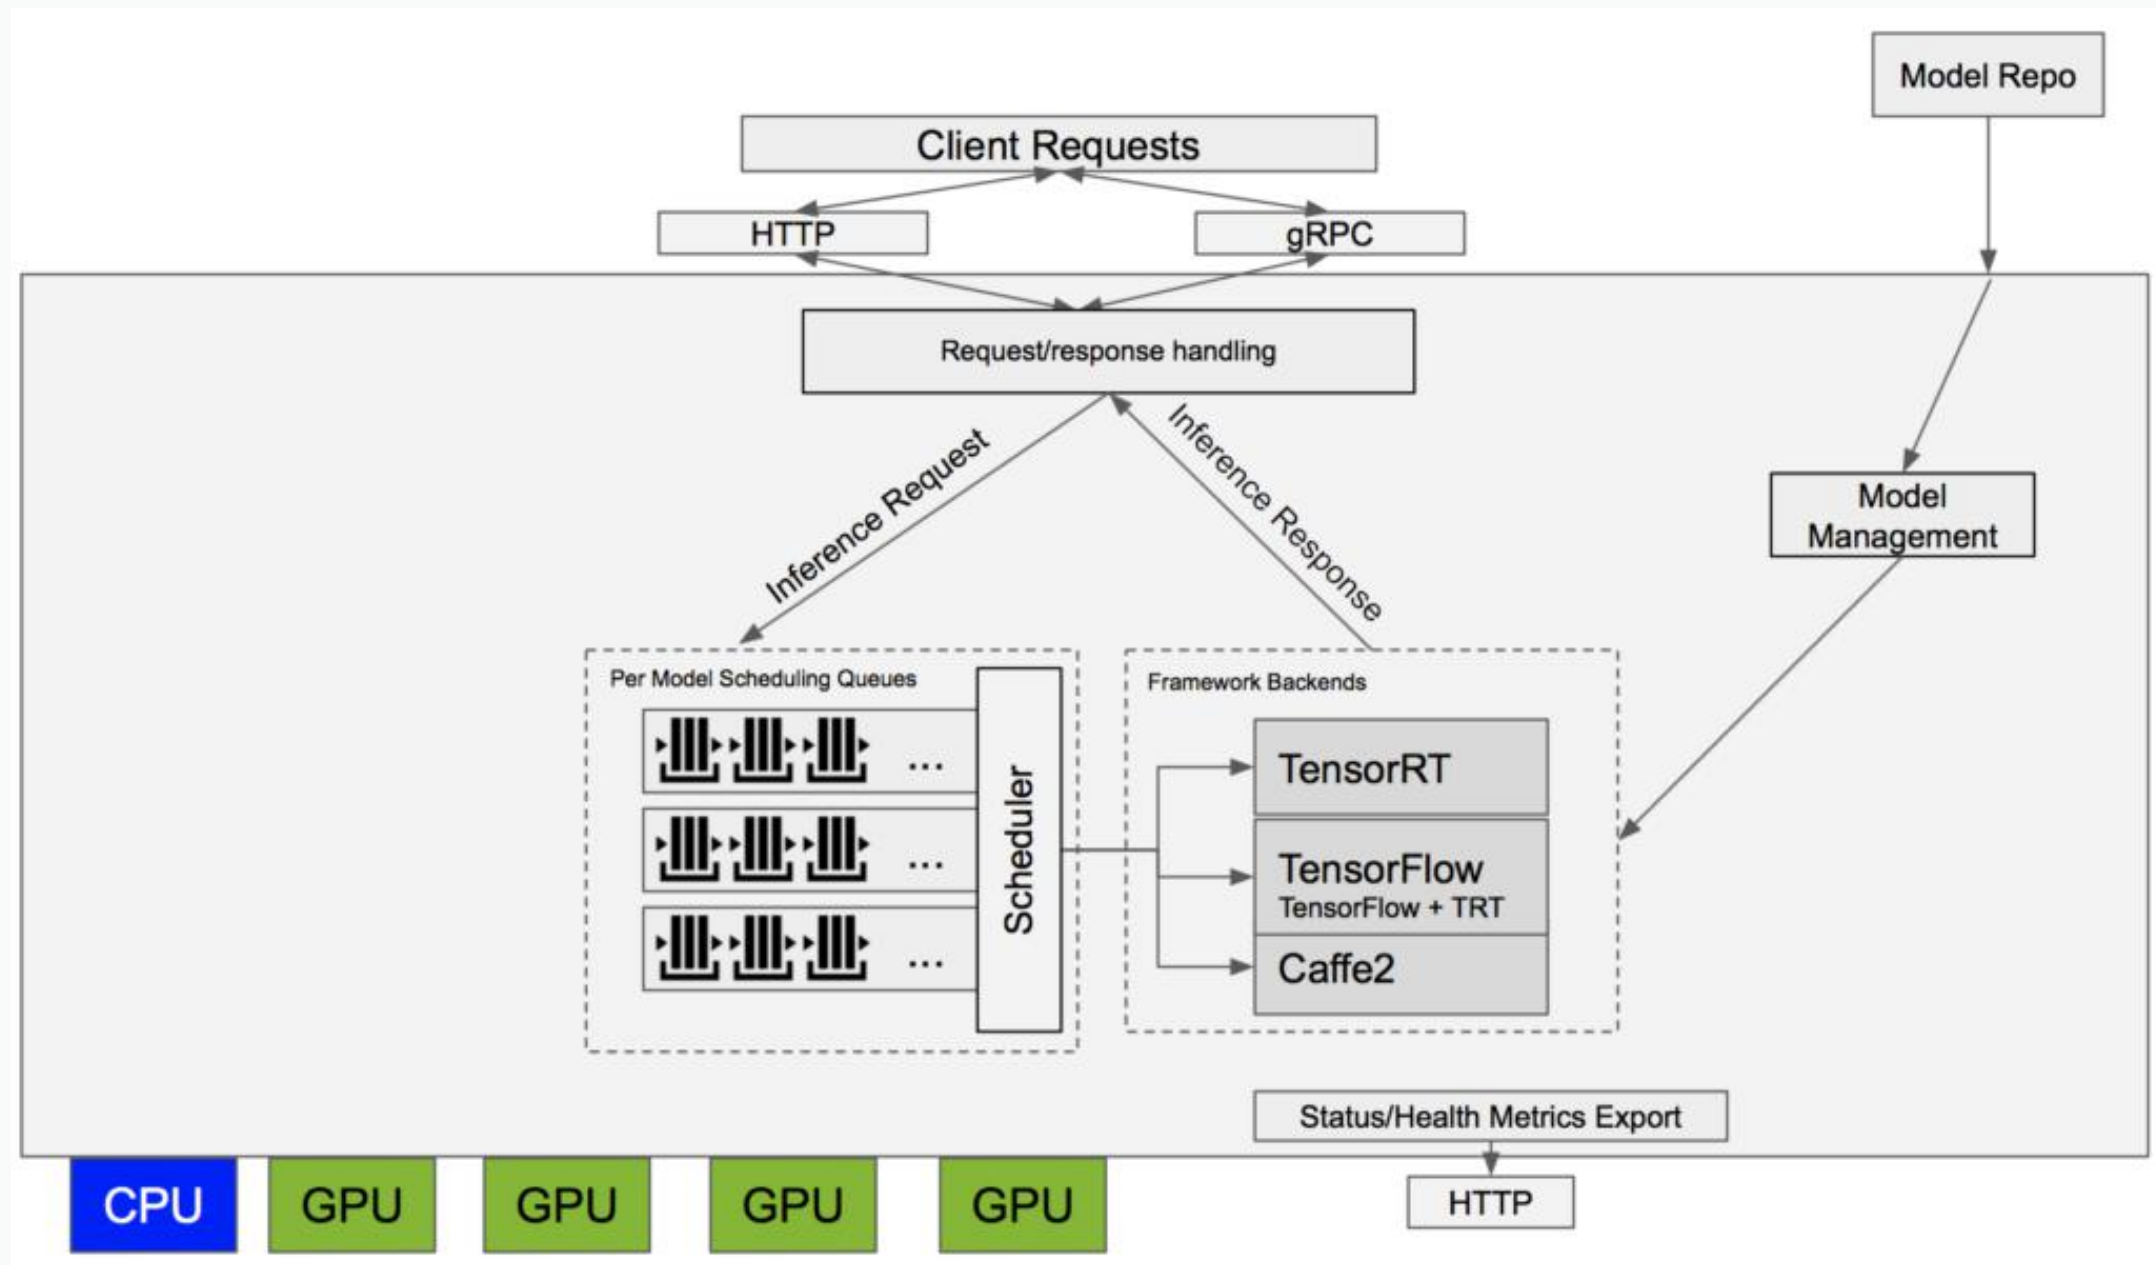
- Min
- Mul
- Neg
- Pad
- ParametricSoftplus
- Pow
- Reciprocal
- ReduceL1
- ReduceL2
- ReduceLogSum
- ReduceLogSumExp
- ReduceMax
- ReduceMean
- ReduceMin
- ReduceProd
- ReduceSum
- ReduceSumSquare
- Relu
- Reshape
- ScaledTanh
- Selu
- Shape
- Sigmoid
- Sin, Cos, Tan, Asin, Acos, Atan, Sinh, Cosh, Asinh, Acosh, and Atanh
- Size
- Slice
- Softmax
- Softplus
- Softsign
- SpaceToDepth
- Split
- Squeeze
- Sub
- Sum
- Tanh
- ThresholdedRelu
- TopK
- Transpose
- Unsqueeze
- Upsample

*Design your network with
tensorRT-supported
layers !*

Demo

(ResNet50)

TensorRT inference server



multiple models

multi-GPUs

multi-framework

model repository

TensorRT inference server

Usage:

```
nvidia-docker run --rm --shm-size=1g --ulimit memlock=-1 \
  --ulimit stack=67108864 \
  -p 8000:8000 -p 8001:8001 -p 8002:8002 \
  -v /path/to/model/repository:/models \
  <tensorrtserver image name> trtserver \
  --model-store=/models
```

Model repository

```
<model-repository-path>/  
  model_0/  
    config.pbtxt  
    output0_labels.txt  
  1/  
    model.plan  
  2/  
    model.plan  
  model_1/  
    config.pbtxt  
    output0_labels.txt  
    output1_labels.txt  
  0/  
    model.graphdef  
  7/  
    model.graphdef
```

Model configuration file

```
name: "mymodel"
platform: "tensorrt_plan"
max_batch_size: 8
input [
  {
    name: "input0"
    data_type: TYPE_FP32
    dims: [ 16 ]
  },
  {
    name: "input1"
    data_type: TYPE_FP32
    dims: [ 16 ]
  }
]
output [
  {
    name: "output0"
    data_type: TYPE_FP32
    dims: [ 16 ]
  }
]
```

Demo

(ResNet50)

References

[1] tensorRT developer guide

https://docs.nvidia.com/deeplearning/sdk/tensorrt-developer-guide/index.html#dla_topic

[2] tensor best practices

<https://docs.nvidia.com/deeplearning/sdk/tensorrt-best-practices/index.html#fusion-types>

[3] tensorRT inference server guide

<https://docs.nvidia.com/deeplearning/sdk/tensorrt-inference-server-guide/docs/index.html>

[4] tensorRT inference server GitHub sources

<https://github.com/NVIDIA/tensorrt-inference-server>

[5] supported layers

<https://docs.nvidia.com/deeplearning/sdk/tensorrt-archived/tensorrt-512rc/tensorrt-support-matrix/index.html>

[6] layer fusion

<https://devblogs.nvidia.com/deploying-deep-learning-nvidia-tensorrt/>

Question ?